

Linear Regression- Vehicle selection

```
In [18]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [3]: #Step-2: Reading the Dataset
df=pd.read_csv(r"C:\Users\sudheer\Downloads\fiat500_VehicleSelection_Dataset (1).csv")
df
```

Out[3]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

```
In [4]: df=df[['age_in_days','price']]
df.columns=['age_in_days','price']
```


```
In [5]: df.head()
```

Out[5]:

	age_in_days	price
0	882	8900
1	1186	8800
2	4658	4200
3	2739	6000
4	3074	5700

```
In [7]: #Step-3: Exploring the Data Scatter - plotting the data scatter
sns.lmplot(x="age_in_days",y="price", data = df, order = 2, ci = None)
df.describe()
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   age_in_days  1538 non-null   int64
1   price       1538 non-null   int64
dtypes: int64(2)
memory usage: 24.2 KB
```



```
In [8]: #Step-4: Data cleaning - Eliminating NaN OR missing input numbers
df.fillna(method = 'ffill', inplace = True)
```

C:\Users\sudheer\AppData\Local\Temp\ipykernel_24120\3221840372.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df.fillna(method = 'ffill', inplace = True)

```
In [10]: # Step-5: Training Our Model
X = np.array(df['age_in_days']).reshape(-1, 1)
y = np.array(df['price']).reshape(-1, 1)
#Seperating the data into independent and dependent variables and convert
#Now each dataset contains only one column
```

```
In [11]: X_train,X_test,y_train,y_test = train_test_split(X, y, test_size = 0.25)
# Splitting the data into training data and test data
regr = LinearRegression()
regr.fit(X_train, y_train)
print(regr.score(X_test, y_test))
```


0.7814006664917448

```
In [12]: #step-6: Exploring Our Results
y_pred = regr.predict(X_test)
plt.scatter(X_test, y_test, color = 'b')
plt.plot(X_test, y_pred,color='k')
plt.show()
# Data scatter of predicted values
```




```
In [13]: # Step-7: Working with a smaller Dataset
df500 = df[:][:500]
# Selecting the 1st 500 rows of the data
sns.lmplot(x = "age_in_days", y = "price", data = df500, order = 1, ci = None)
```

Out[13]: <seaborn.axisgrid.FacetGrid at 0x1d97fb8e8c0>



```
In [15]: df500.fillna(method = 'ffill', inplace = True)
X = np.array(df500['age_in_days']).reshape(-1, 1)
y = np.array(df500['price']).reshape(-1, 1)
df500.dropna(inplace = True)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25)
regr = LinearRegression()
regr.fit(X_train, y_train)
print("Regression:", regr.score(X_test, y_test))
y_pred = regr.predict(X_test)
plt.scatter(X_test, y_test, color = 'b')
plt.scatter(X_test, y_pred, color = 'k')
plt.show()
```

Regression: 0.8195780898732068



```
In [16]: #Step-8: Evaluation of model

from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

#Train the model

model = LinearRegression()
model.fit(X_train, y_train)

#Evaluating the model on the test set

y_pred = model.predict(X_test)

r2 = r2_score(y_test, y_pred)

print("R2 score:", r2)
```

R2 score: 0.8195780898732068

Step-9:Conclusion:

Dataset we have taken is poor for Linear Model, but with the smaller data works well with Linear Model.

In []: